B Inside temporal processing

Shown: three frames with a 1×1 convolution at each resolution. Each row has its own trainable weights.

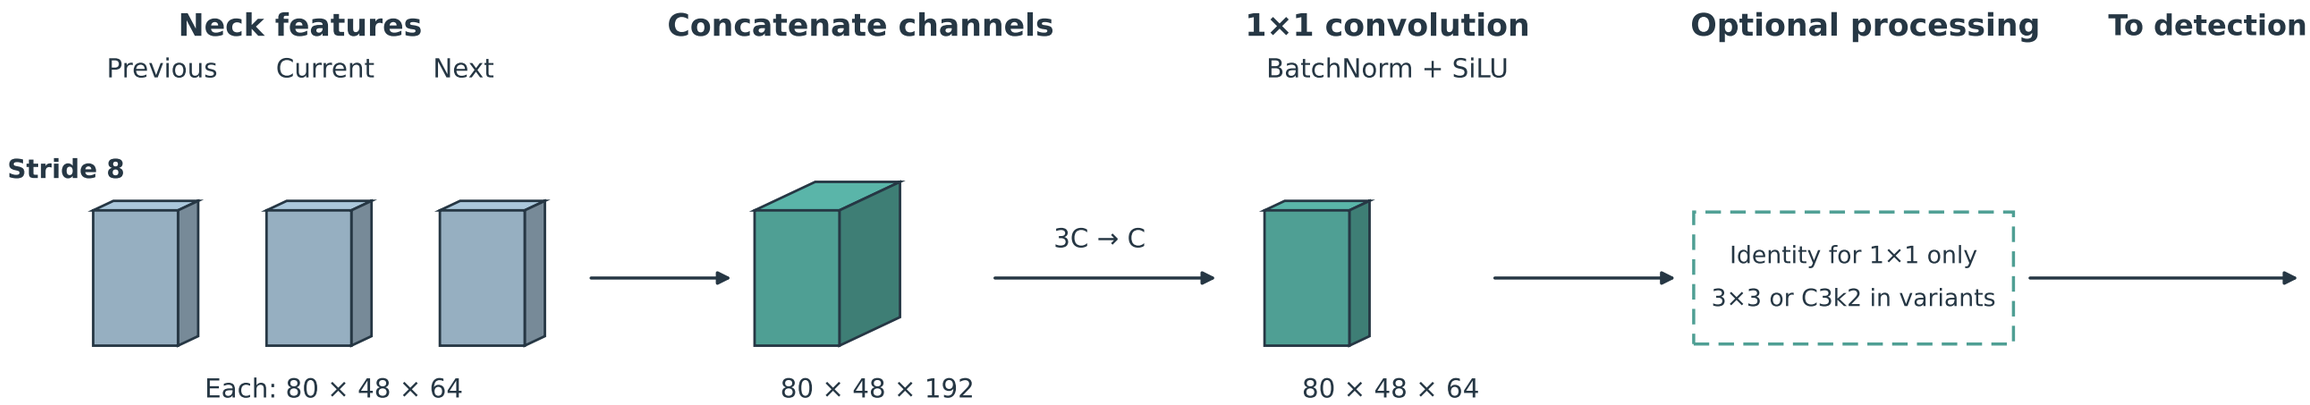

Dimensions are width × height × channels for a 640 × 384 input. Batch dimension omitted; block sizes are schematic.

Repeated frame controls: all three inputs come from t. Five frame variants concatenate t − 2 through t + 2 (5C → C).

The five frame spatial block is 3×3 or 5×5. Outputs preserve each resolution and its channel count.